

quantmsdiann: a scalable SDRF-driven DIA-NN workflow for reanalysis of single-cell, spatial, and bulk proteomics datasets

Qi-Xuan Yue^a, Yufei Shen^a, Asier Larrea^c, Chengxin Dai^b, Henry Webel^d, Jose Nimo^f, Fabian Coscia^f, Orhun Kok^g, Nikolai Slavov^{g,h}, Mingze Bai^a, Vadim Demichev^e, Yasset Perez-Riverol^{c,*}

^a*Chongqing Key Laboratory of Big Data for Bio Intelligence, Chongqing University of Posts and Telecommunications, Chongqing, China*

^b*State Key Laboratory of Proteomics, Beijing Proteome Research Center, National Center for Protein Sciences (Beijing), Beijing Institute of Life Omics, Beijing, China*

^c*European Molecular Biology Laboratory, European Bioinformatics Institute (EMBL-EBI), Wellcome Genome Campus, Hinxton, Cambridge, UK*

^d*Novo Nordisk Foundation Center for Protein Research, Faculty of Health and Medical Sciences, University of Copenhagen, Copenhagen, Denmark*

^e*Quantitative Proteomics laboratory, Charité – Universitätsmedizin Berlin, Berlin, Germany*

^f*Spatial Proteomics Group, Max-Delbrück-Center for Molecular Medicine in the Helmholtz Association (MDC), Berlin, Germany*

^g*Department of Bioengineering, Department of Biology, Department of Chemistry and Chemical Biology, Single Cell Proteomics Center, and Barnett Institute for Chemical and Biological Analysis, Northeastern University, Boston, Massachusetts, USA*

^h*Parallel Squared Technology Institute, Watertown, Massachusetts, USA*

Abstract

Public proteomics archives now host thousands of data-independent acquisition (DIA) datasets, but they remain analytically siloed because each was processed with a different software configuration. We present **quantmsdiann**, an open-source Nextflow/nf-core workflow that parallelises DIA-NN across cloud and HPC infrastructures, guided by the experimental design declared in SDRF format. The workflow ships container-pinned profiles for any DIA-NN version through Docker and Singularity, requiring no manual installation or dependency resolution, and natively handles all vendor formats, converting to

*Corresponding author.

URL: yasset@ebi.ac.uk (Yasset Perez-Riverol)

mzML when needed (e.g. SCIEX). It exports harmonised quantification tables as MSstats input, the new Quantitative Proteomics eXchange format (QPX), and a `pmultiqc` quality-control report. The parallelisation design allows a 2,300-run single-cell dataset to be reanalysed in 2.4 hours on an HPC cluster of 300 nodes.

We benchmark `quantmsdiann` on the full ProteoBench DIA panel (five DIA-NN versions $\times$ four instrument modalities), reanalyse three independent bulk cell-line cohorts (PXD003539, PXD004701, PXD030304), and assemble a pan-cohort of five public DIA cell-line reanalyses under a single SDRF-driven invocation; SDRF-driven reprocessing with current DIA-NN versions matches or exceeds the originally deposited results. Single-cell and spatial proteomics demonstrations are part of ongoing work. `quantmsdiann` is available at <https://github.com/bigbio/quantmsdiann>.

Keywords: DIA-NN, single-cell proteomics, spatial proteomics, plexDIA, data-independent acquisition, Nextflow, nf-core, SDRF, reproducibility, QPX, `pmultiqc`

1. Introduction

Data-independent acquisition (DIA) is now the dominant mass-spectrometry strategy for proteomics at scale, and DIA-NN [1] is among the most widely adopted analysis engines across timsTOF [2, 3], Astral, and Orbitrap platforms. Recent benchmarks confirm DIA-NN’s competitiveness for single-cell DIA workflows [4], and versions 1.9.x and 2.x have expanded support for plexDIA, timsTOF parallel accumulation-serial fragmentation (PASEF), and vendor-native raw-file reading. Despite this analytical maturity, most DIA-NN users run the tool as a desktop application, which limits its scalability for large datasets. Moreover, deployments on HPC and cloud infrastructure remain artisanal, relying on shell scripts, ad-hoc directory conventions, and hand-curated sample sheets, so the same raw data analysed by two laboratories, or with two DIA-NN releases, can yield non-overlapping quantification tables that are difficult to reconcile.

Beyond per-study analysis, the value of public proteomics archives such as PRIDE Archive [5] lies as much in systematic reanalysis as in original deposition. In 2024 we published bigbio/quantms [6], a Nextflow pipeline that standardised multi-engine DIA and DDA analyses behind a Sample and Data Relationship Format (SDRF) [7] input on the nf-core framework [8]. `quantms`

was, however, built around DIA-NN 1.8.1 with a fixed parameter set, and its DDA/DIA breadth made extending DIA-specific functionality difficult: new DIA-NN parameters conflicted with DDA options, versions were hard-coded, sample-level acquisition parameters were not properly propagated to per-run DIA-NN invocations, and raw-file conversion was mandatory. Iterative reanalysis at archive scale requires a workflow that combines the reproducibility of nf-core, the sensitivity of current DIA-NN releases, the metadata discipline of SDRF, and a parallelisation strategy capable of processing thousands of raw files.

Here, we present **quantmsdiann**, an independent SDRF-driven Nextflow pipeline dedicated to DIA-NN that emerged from the need to extend DIA-specific features beyond what the original bigbio/quantms pipeline could accommodate. While **quantmsdiann** shares its SDRF-first design philosophy and parts of the configuration-generation tooling with quantms, its workflow, container set, and module library are developed and released independently, with container-pinned profiles for every DIA-NN version selectable at runtime under Docker or Singularity. Downstream, it exports a harmonised QPX (Quantitative Proteomics eXchange) Parquet archive alongside the standard DIA-NN reports and MSstats input, bundling the SDRF, identifications, quantifications, and pipeline provenance into a single queryable artefact for reanalysis and long-term archival. We demonstrate the pipeline along three axes: ProteoBench [9] benchmarks across four DIA modalities and five DIA-NN releases; three independent bulk-DIA cell-line reanalyses (PXD003539, PXD004701, PXD030304); and a pan-cohort of five public DIA cell-line reanalyses integrated under a single SDRF-driven invocation.

2. Experimental Procedures

2.1. Pipeline implementation

quantmsdiann is implemented in Nextflow DSL2 ($\geq 25.10.4$) and follows nf-core community standards [8] for pipeline structure, testing, documentation, and continuous integration. Each computational step is packaged as a versioned container image (Docker, Singularity, or Apptainer); upstream community tools (DIA-NN 1.8.1, ThermoRawFileParser [10], pridepy [11], qpx) are pulled from BioContainers, while pipeline-specific images (newer DIA-NN releases and **wiffconverter**) are hosted on ghcr.io/bigbio. The repository at <https://github.com/bigbio/quantmsdiann> is organised as a standard nf-core pipeline with main workflow (**main.nf**), subworkflows, modules, test

data and CI configurations, and documentation. Pipeline version v2.1.0 (release codename “Sao Pablo”, 2026-05-08) was used for the analyses reported in this paper.

2.2. SDRF parsing and configuration generation

`quantmsdiann` accepts SDRF [7] as primary input. SDRF parsing is performed using the `sdrf-pipelines` Python package developed within the `quantms` project, which validates the SDRF against the proteomics-specific subset of EFO, PSI-MS, and PRIDE ontology terms and emits a normalised per-sample channel for downstream Nextflow processes. `DIA-NN` configuration generation extracts enzyme, fixed and variable modification, instrument, precursor mass tolerance, fragment mass tolerance, and missed-cleavage settings from the SDRF and writes a `DIA-NN` command-line configuration. Where SDRF columns are absent, `DIA-NN` defaults appropriate for the inferred instrument are applied.

2.3. Raw file handling and DIA-NN execution

Raw inputs are dispatched to format-specific handlers in the file-preparation subworkflow, each conversion or decompression running as an independent, parallelised Nextflow process. Thermo `.raw` files are read natively by `DIA-NN` $\geq 2.1.0$ (which added Linux native-reader support) or converted to `mzML` with `ThermoRawFileParser` [10] otherwise (`-mzml_convert`); Bruker `.d` folders are ingested natively, with compressed forms (`.d.tar/.zip/.gz`) decompressed first; and Sciex `.wiff` / `.wiff2` are converted to `mzML` with the open-source `wiffconverter` (<https://github.com/bigbio/wiffconverter>), removing the historical dependency on a proprietary ProteoWizard installation. Generic `.dia` and pre-converted `.mzML` inputs pass through unchanged, with optional OpenMS re-indexing (`-reindex_mzml`). `DIA-NN` [1] is then invoked at the version chosen via `-diann_version`: the ProteoBench analyses use `DIA-NN` 1.8.1, `DIA-NN` 2.1.0, `DIA-NN` 2.2.0, `DIA-NN` 2.3.2, and `DIA-NN` 2.5.0 from the same SDRF, and the cell-line reanalyses use `DIA-NN` 2.5.0 unless otherwise stated.

2.4. Spectral-library generation

`quantmsdiann` runs `DIA-NN` library-free, generating the spectral library in silico from the input FASTA rather than supplying an external experimental library: `DIA-NN`’s deep-learning predictor builds a predicted library from the protein sequences, each run is searched against it, a cohort-wide empirical library is assembled from the confident first-pass identifications, and a final pass

re-searches and quantifies against that refined library. The source FASTA was the UniProt human reference proteome [12] (UP000005640, *[release]*), with match-between-runs enabled at the cohort level. In ProteoBench’s submission taxonomy this is the predicted-library strategy (`predictors_library = DI-ANN`); the ProteoBench comparison (Sec. 3.3) is therefore restricted to community submissions that used the same predicted-from-FASTA strategy.

2.5. *Quality control and QPX export*

Per-run and per-cohort quality control reports are generated by `pmultiqc` [13], a proteomics extension of MultiQC [14] that summarises DIA-NN-specific metrics including peptide and protein counts, missed cleavages, charge-state distributions, precursor and fragment mass accuracy, and per-run identification yield. The QPX archive is produced by the `qpxc` CLI from the `qpx` package (<https://github.com/bigbio/qpx>), which writes Parquet files for identifications (psm, peptide, protein-group, feature), metadata (sample, run, ontology), and provenance, and optionally exports MuData (`.h5mu`) containers for scverse-based [15] downstream analysis.

2.6. *Datasets benchmarked*

The benchmark and reanalysis datasets used in this study (ProteoBench DIA modules 5/7/9/10, the three per-cohort cell-line reanalyses PXD003539, PXD004701, PXD030304, the two additional pan-cohort datasets PXD041421 and PXD017199, and the PXD071075 scalability sweep) are summarised together with their instrument, acquisition mode, and deposition year in Additional file 1: Table S1. For each dataset, an SDRF was prepared from the original sample metadata and is provided as part of the Additional file 1 deposition. For each cell-line reanalysis cohort, the original published quantification was downloaded from the PRIDE deposition and used as the comparator, thresholded at 1% precursor-level FDR for both `quantmsdiann` and the original analyses. The conservative contaminant/decoy filter and the like-for-like-threshold rule applied to the original-paper headline counts are described in Additional file 1.

2.7. *Compute environment and benchmarking*

Runtime benchmarks were performed on the EMBL-EBI HPC cluster (LSF scheduler), with per-task CPU and memory allocations dispatched by Nextflow. Memory profiling was performed using the Nextflow trace report (`-with-trace`).

3. Results

3.1. End-to-end throughput on archive-scale DIA cohorts

The headline operational result of `quantmsdiann` is that a single SDRF-driven invocation processes archive-scale DIA cohorts within hours of wall-clock time on commodity HPC, with wall-clock time decreasing monotonically with cluster width over the tested range up to a practical knee near 200 nodes. On the PXD071075 single-cell DIA benchmark ($\sim 2,300$ raw files, Wang *et al.* [4]), increasing the Nextflow `executor.queueSize` from 10 to 300 cluster nodes reduced end-to-end wall-clock time from 37.7 h to 2.4 h (Fig. 1b). Scaling was monotonic across the sweep, with a practical knee near 200 nodes (2.6 h) beyond which the cohort-level consolidation step dominates and additional parallel workers yield diminishing returns. Beyond the 200-node knee, gains in wall-clock time come at decreasing CPU-hour efficiency: scaling from 10 to 300 nodes reduces wall-clock by $15.5\times$ for a $30\times$ increase in cluster width, so users should choose cluster width based on the desired turn-around time rather than the maximum-throughput operating point. Across ProteoBench benchmark datasets and our public cell-line cohorts, final-run wall-clock time tracked cohort size and instrument family without per-dataset workflow changes (Fig. 1c): the largest cohort in the panel, PXD030304 (5,798 TripleTOF 6600 runs), completed in 4.8 h on the same configuration that finishes the 6-run ProteoBench cohorts in under one hour. Peak memory usage during the consolidation step remained below $\lceil GB \rceil$ GB for cohorts up to $\lceil N_{max\ raw\ files} \rceil$ raw files, allowing `quantmsdiann` to be deployed on commodity 64–128 GB HPC nodes rather than requiring high-memory specialist hardware. Per-step wall-clock and resource envelopes are reported in Additional file 1: Figs. S1 and S2.

3.2. `quantmsdiann` pipeline architecture

`quantmsdiann` is implemented in Nextflow DSL2 [8] and follows nf-core community standards, distributing each computational step as a versioned Docker or Singularity/Apptainer container. The pipeline accepts a single SDRF file as primary input, which encodes per-sample metadata (cell line or tissue origin, factor values, instrument, post-translational modifications) and per-run file references. From this input, `quantmsdiann` derives the analysis parameters required by DIA-NN automatically, eliminating the most common source of analysis-script divergence between research groups. The workflow comprises four mandatory and three optional modules organised around the

DIA-NN analysis core (Fig. 1a), alternating between per-file parallel stages (red; in-silico library generation when requested, preliminary calibration, individual search) and collective stages (green; empirical library assembly, final quantification, MSstats conversion, `pmultiqc` [13] quality-control reporting) operating on the full cohort.

3.2.1. Parallelisation strategy

Raw-file conversion and per-run quantification are dispatched as independent Nextflow tasks, enabling the throughput scaling reported in Section 3.1. Dataset-wide DIA-NN steps that require the full cohort (match-between-runs, normalisation, protein inference) run in a single consolidation task with tuned memory ceilings, avoiding out-of-memory failures common to single-machine DIA-NN runs on 1,000+ file inputs. Outputs include the standard DIA-NN report tables (`report.tsv`, `report.pr_matrix.tsv`, `report.pg_matrix.tsv`) alongside MSstats-formatted input and a `pmultiqc` [13] quality-control report. A harmonised QPX archive bundles the SDRF, identifications (PSM, peptide, protein-group, feature), MS metadata, ontology terms, and pipeline provenance into a Parquet-based artefact, which can be queried directly with DuckDB and, for the quantification layers, exported to MuData (`.h5mu`) for scverse-compatible [15] downstream analysis.

3.2.2. DIA-NN version flexibility

Container-pinned profiles for DIA-NN 1.8.1, 2.1.0, 2.2.0, 2.3.2 and 2.5.0 are shipped with the pipeline, with the DIA-NN binary selected at runtime through the `-diann_version` parameter; users can add further versions by supplying a custom container. Identical SDRF inputs can therefore be processed against multiple DIA-NN versions in one invocation, with separate output directories per version. This explicit version handling enables cross-version reproducibility audits without re-engineering the workflow and supports use cases such as reanalysis of historical datasets with the DIA-NN version originally used by the depositing laboratory, alongside re-quantification with the current release.

3.2.3. Library handling

For single-cell and spatial DIA experiments, matched empirical spectral libraries from bulk samples are often unavailable. `quantmsdiann` supports two complementary library modes: (i) DIA-NN library-free analysis from a user-supplied FASTA proteome (default), in which DIA-NN generates a predicted in-silico library and a cohort-wide empirical library is then assembled across all

runs in the SDRF; and (ii) user-supplied spectral library passed via `-speclib`, for users wishing to reuse a previously generated empirical or hybrid library. In-silico library generation, per-run calibration, and empirical-library assembly are each parallelised across the relevant runs.

3.3. *ProteoBench benchmark across DIA-NN versions and instruments*

To demonstrate that `quantmsdiann` preserves DIA-NN’s analytical performance under SDRF-driven orchestration, we processed every public ProteoBench DIA module through the workflow and compared the resulting identification counts against the ProteoBench public-submission snapshot (29 May 2026) [9]. `quantmsdiann` runs DIA-NN library-free: it predicts the spectral library in-silico from the FASTA and refines it from the acquired data, with no externally supplied experimental library. For a like-for-like comparison we therefore restricted the community set to ProteoBench DIA-NN submissions that used the same predicted-from-FASTA library (`predictors_library = DIA-NN`); submissions that loaded a pre-existing experimental or user-curated spectral library search a different candidate space and are excluded from the headline comparison. The benchmark panel covers four instrument families and acquisition modes: Orbitrap Astral 2 Th DIA (*Module 7*), single-cell DIA on Orbitrap (*Module 9*, PXD049412), timsTOF diaPASEF (*Module 5*, PXD062685), and SCIEX ZenoTOF 7600 (*Module 10*, PXD070049). For each module we ran the workflow against five DIA-NN releases (1.8.1, 2.1.0, 2.2.0, 2.3.2, and 2.5.0) in a single invocation, with all five versions parametrised from the same SDRF (Fig. 2).

Precursor counts (quantified in at least one replicate) grew monotonically with DIA-NN version on every module, rising by 14% on Module 7 (105,793 $\rightarrow$ 120,973 from 1.8.1 to 2.5.0), 14% on Module 9 (38,218 $\rightarrow$ 43,761), 27% on Module 5 (91,075 $\rightarrow$ 115,735), and 17% on Module 10 (85,847 $\rightarrow$ 100,524). Protein-group counts at 1% FDR increased over the same range (endpoint gains of 7% / 8% / 14% / 7% respectively), though not strictly monotonically: on each module the protein count dips by 1–2% at an intermediate version before the newest release recovers the headline gain. Reproducibility-filtered precursor counts (≥ 3 replicates) can likewise be non-monotonic — on the ZenoTOF module (PXD070049), for instance, DIA-NN 2.5.0 yields the most total precursors but slightly fewer reproducible ones than 2.3.2. This monotonic version-progress trend is consistent with what the broader ProteoBench community reports across DIA-NN releases [9], confirming that the SDRF wrapper introduces no version-coupling artefacts. Quantification accuracy, by contrast,

is essentially invariant across DIA-NN versions. On every module the measured per-species $\log_2$ fold-changes recover the expected HYE ratios—human at unity, yeast and *E. coli* near their design ratios, with the mild attenuation at the extreme *E. coli* ratio that is characteristic of DIA quantification—and the five DIA-NN releases overlie one another at every species (Fig. 2b). The version-to-version shift in median quantification error $|\epsilon|$ ($\sim 0.01 \log_2$, $\approx 0.5\text{--}1\%$ in linear fold-change) is far smaller than the per-precursor spread, so the choice of DIA-NN release is best driven by the identification gains in panel (a) rather than by accuracy. On the two modules for which the snapshot contained predicted-library DIA-NN submissions—Orbitrap Astral (Module 7, $n = 15$) and timsTOF diaPASEF (Module 5, $n = 8$)—**quantmsdiann**’s five versions fall within the community median- $|\epsilon|$ distribution, at or below its median (Fig. 2c); the single-cell and ZenoTOF modules carried no predicted-library community submission. Parameter-matched per-submission comparisons are provided in Additional file 1: Fig. S3.

3.4. *Independent cell-line reanalyses recover comparable or deeper coverage than the original quantifications*

To evaluate **quantmsdiann** on production-scale public deposits, we reanalysed three independent bulk-DIA cell-line cohorts that had been processed and published with different DIA engines (Fig. 3). For each cohort the same raw files were retrieved from PRIDE Archive [5], parsed through an SDRF generated from the original sample annotations, and run through **quantmsdiann** with current DIA-NN releases.

On the Guo *et al.*, 2019 NCI-60 panel (PXD003539; originally analysed with OpenSWATH using the Guo 2019 assay library), **quantmsdiann** recovered 95,633 peptides at 1% FDR against 40,592 in the original submission (+135%), with comparable protein-group counts (6,747 versus 6,556; Fig. 3a). Because the original Guo 2019 analysis used OpenSWATH with a fixed assay library while **quantmsdiann** runs DIA-NN in library-free mode, this gain reflects both the SDRF-driven orchestration and the change of analysis engine; a like-for-like DIA-NN-on-original-library control is part of ongoing work. On the Sun *et al.*, 2023 breast-cancer panel of 76 cell lines (PXD004701; originally analysed with PCT-SWATH), **quantmsdiann** recovered 6,146 protein groups under the original consistency filter (proteotypic peptides, $\leq 90\%$ missing) against 6,091 in the published analysis (+0.9%), with raw-matrix counts within 10% on the unfiltered metrics (Fig. 3b). On the ProCan-DepMapSanger 949-cell-line pan-cancer panel (PXD030304; originally

analysed with DIA-NN 1.7.x), `quantmsdiann` recovered 8,921 protein groups at 1% FDR against 8,498 originally (+5%) and 7,990 protein groups requiring ≥ 2 unique peptides versus 6,692 (+19%; Fig. 3c). Both the original 8,498 protein-group headline (Global.Q.Value ≤ 0.01 , no per-protein peptide threshold) and the stringent 6,692 count (same Global.Q.Value filter plus a ≥ 2 supporting-peptide requirement) are reported in the ProCan-DepMapSanger paper itself; the `quantmsdiann` numbers (8,921 and 7,990) are computed under the matching filters and the same conservative contaminant/decoy filter, so each side of the comparison is like-for-like. Across the three cohorts, the gain was largest on the deposits that had been processed with older or non-DIA-NN engines (Guo 2019 OpenSWATH, ProCan 1.7.x); the cohort that had already been analysed with a current DIA engine (Sun 2023) yielded near-parity, as expected. Per-condition completeness, peptide-per-protein distributions, gene-versus-accession overlap and per-cell-line missingness for each cohort are provided as Additional file 1: Figs. S5–S10.

3.5. Pan-cohort of five public DIA cell-line reanalyses

To illustrate the value of running the same SDRF-driven workflow across many independent deposits, we combined the three reanalysed cohorts above with two additional cell-line DIA panels (PXD041421, Wang 2023; PXD017199, Tognetti 2021, 67 breast cancer cell lines) into a single `quantmsdiann` invocation and assembled a pan-cohort view of per-cohort yield and per-tissue coverage (Fig. 4). Per-cohort headline counts (Fig. 4a) confirm that `quantmsdiann` matches or exceeds the original published count on each cohort. Across the five cohorts the reanalysis spans 12,229 unique UniProt accessions, of which 5,144 (42.1%) are detected in all five—a pan-cohort core proteome—and a further 1,334 (10.9%) in exactly four; these overlap counts are tabulated in the combined per-cohort counts provided with the analysis code.

Tissue-level coverage across the combined panel spans 27 Cellosaurus-derived tissue categories (plus the catch-all “Other tissue” and “Healthy (Non-cancer)” bins; Fig. 4b,c): lung, breast, haematopoietic-and-lymphoid, skin, and central nervous system together account for 59% of all cell-line entries in the pan-cohort. Because every count comes from running a single SDRF-parameterised `quantmsdiann` invocation rather than reconciling per-study quantifications, the resulting protein matrix is directly queryable from the published QPX archive and can be filtered, joined, or re-grouped by Cellosaurus, NCIt, or Disease Ontology terms without manual harmonisation.

The consolidated SDRF and per-cohort `quantmsdiann` configurations underpinning this pan-cohort run will be deposited on Zenodo before publication, so that the full input–configuration–output triple can be re-executed by third parties.

4. Discussion

`quantmsdiann` lowers the engineering barrier between a deposited DIA dataset and a reproducible quantification table, and it does so at a scale that makes systematic reanalysis of public DIA archives tractable for the first time. The SDRF-first design encourages experiment annotation upstream of analysis and produces metadata that is interoperable with the wider proteomics ecosystem [7]; the QPX archive bundles the SDRF, harmonised quantification tables and pipeline provenance into a single Parquet container so that reanalysed results can be archived (e.g. on Zenodo) and re-queried alongside the original raw data. By focusing exclusively on DIA-NN, rather than the multi-engine scope of the related bigbio/quantms pipeline [6], `quantmsdiann` achieves a leaner dependency surface and faster pipeline-level execution for the DIA-only regime, where the dominant cost is per-file DIA-NN invocation rather than across-engine comparison.

The demonstrations we present stress the workflow along three orthogonal axes. The ProteoBench benchmark (Fig. 2) tests analytical fidelity across four instrument modalities and five DIA-NN releases; the three independent cell-line reanalyses (Fig. 3) test recovery against published per-study quantifications, with results ranging from near-parity (Sun 2023, current DIA engine) to substantial recovery gains (Guo 2019 OpenSWATH; ProCan DIA-NN 1.7.x); and the pan-cohort of five public DIA cell-line reanalyses (Fig. 4) tests archive-scale throughput on five concurrently reanalysed deposits. Single-cell and spatial proteomics biology demonstrations are part of ongoing work and are not part of this submission’s biology figures; the single-cell PXD071075 cohort used as the throughput-scaling benchmark in Fig. 1 is, however, real production data. This pattern argues that the value of DIA reanalysis is real but time-limited—largest for deposits originally analysed with older or non-DIA-NN engines and shrinking toward parity for recently processed cohorts—exactly the regime an SDRF-driven workflow is well-positioned to address as new DIA-NN versions ship. Demonstrations on flagship single-cell DIA and laser-microdissected spatial cohorts are part of the development roadmap; the workflow already supports these acquisition modes (Fig. 1a,




Figure 1: **quantmsdiann pipeline architecture and runtime scaling.** (a) Subway diagram of the **quantmsdiann** workflow. Mandatory modules are shown as filled circles, optional/skippable modules as open circles; per-file parallel stages are marked in red, collective cohort stages in green. The pipeline consumes an SDRF and a FASTA reference, optionally converts vendor raw files to mzML, generates an in-silico or empirical spectral library, runs per-file DIA-NN analyses in parallel, consolidates results in a single cohort-level final quantification, and emits standard DIA-NN tables, an MSstats input, a QPX Parquet archive, and a **pmultiqc** quality-control report. (b) Workflow wall-clock time versus Nextflow `executor.queueSize` (log-scaled cluster nodes; ticks at 10, 50, 100, 200, 300) on the PXD071075 single-cell cohort (Wang *et al.*, 2025); wall-clock drops from 37.7 h at 10 nodes to 2.4 h at 300 nodes, with the practical knee near 200 nodes. (c) Final-run wall-clock versus number of MS data files across ProteoBench benchmark and cell-line cohorts, coloured by instrument family; the PXD071075 production run (300 nodes) is annotated.




Figure 2: quantmsdiann on ProteoBench across five DIA-NN releases and four DIA modalities (Orbitrap Astral, Module 7; single-cell, Module 9/PXD049412; timsTOF diaPASEF, Module 5/PXD062685; SCIEX ZenoTOF, Module 10/PXD070049). (a) Precursors at 1% FDR by DIA-NN version (1.8.1–2.5.0; light to dark blue) from one quantmsdiann invocation; precursor counts grow monotonically with version (e.g. Module 5 91,075 → 115,735). (b) Per-species quantification accuracy: measured $\log_2$ fold-change for each HYE species against the ProteoBench-expected ratio (dashed), with all five DIA-NN versions overlaid (light to dark). The versions cluster tightly on every species and module — quantification accuracy is essentially invariant across DIA-NN releases, the version-to-version median $|\epsilon|$ shift ($\sim 0.01 \log_2$) being far smaller than the per-precursor spread; the mild attenuation at the extreme *E. coli* ratio is the expected DIA ratio-compression. (c) For the two modules with predicted-from-FASTA (`predictors_library` = DIANN) DIA-NN community submissions (snapshot 29 May 2026), quantmsdiann’s five versions (red) sit within the community median- $|\epsilon|$ distribution (grey; box = IQR). The single-cell and ZenoTOF modules had no predicted-library community submission. See Additional file 1: Fig. S3 for parameter-matched per-submission comparisons.




Figure 3: Independent reanalysis of public DIA datasets, spanning bulk cell lines and single cells, recovers as many or more identifications than the originally deposited quantifications. Each panel compares the originally published peptide and/or protein-group counts (grey) against the **quantmsdiann** reanalysis at 1% FDR (blue) using the same raw data. **(a)** Guo *et al.*, 2019 NCI-60 OpenSWATH analysis (PXD003539, original engine = OpenSWATH with the Guo 2019 assay library): **quantmsdiann** recovers 95,633 peptides versus 40,592 originally (+135%) and a comparable 6,747 protein groups versus 6,556. Because **quantmsdiann** runs DIA-NN in library-free mode while the original ran OpenSWATH against a fixed assay library, this gain reflects both SDRF-driven orchestration and the change of analysis engine; a like-for-like DIA-NN-on-original-library control is part of ongoing work. **(b)** Galatidou *et al.*, 2024 single-cell oocyte plexDIA (MSV000093870; mTRAQ 3-channel, Q Exactive): from the same raw data **quantmsdiann** quantifies 2,904 protein groups versus 2,122 in the original analysis (+37%), recovering 95% of the original accessions plus 1,539 additional proteins at comparable per-cell median depth (1,736 versus 1,784 protein groups; full per-cell, total and accession-overlap comparison in Additional file 1). This is the first plexDIA cohort processed through **quantmsdiann**. **(c)** ProCan-DepMapSanger 2022 pan-cancer cell-line panel (PXD030304, 949 lines): **quantmsdiann** recovers 8,921 protein groups at 1% FDR versus 8,498 originally (+5%), and 7,990 versus 6,692 after requiring ≥ 2 unique peptides per protein (+19%). Both the original 8,498 headline (Global.Q.Value ≤ 0.01) and the stringent 6,692 count (Global.Q.Value ≤ 0.01 plus ≥ 2 supporting peptides) are reported in the ProCan-DepMapSanger paper itself; the **quantmsdiann** counts are computed under the matching filters, so each side of the comparison is like-for-like. Per-condition completeness, peptide-per-protein distributions and gene- versus accession-level overlap are shown in Additional file 1: Figs. S5–S10.




Figure 4: **Pan-cohort of five public DIA cell-line reanalyses by quantmsdiann.** Combines PXD003539 (NCI-60, Guo 2019), PXD030304 (ProCan-DepMapSanger, 949 lines), PXD004701 (Sun 2023, 76 breast cancer lines), PXD041421 (Wang 2023), and PXD017199 (Tognetti 2021, 67 breast cancer lines), reanalysed under a single **quantmsdiann** invocation. (a) Per-cohort headline protein-group counts: the originally published count (grey) and the **quantmsdiann** reanalysis (blue); PXD017199 and PXD041421 have no published DIA headline, so only the **quantmsdiann** bar is shown. (b) Cell lines per pan-cancer tissue, stacked by contributing cohort. (c) Unique UniProt proteins detected per tissue across the combined panel.

library-handling branch) and SDRFs for the Brunner [2], Derks [16], Leduc [17], Wang [4], Mund [18] and Makhmut [19] datasets are in preparation.

The cross-DIA-NN-version reproducibility view enabled by `quantmsdiann` is, to our knowledge, the first pipeline-level treatment of this question for DIA-NN. We expect this to become increasingly important as DIA-NN evolves: a deposited single-cell DIA dataset analysed with DIA-NN 1.8 today and re-analysed with DIA-NN 2.x in two years should not produce different biological conclusions for the same data and parameters, but in current practice such re-analyses are rarely performed because the bespoke per-study scripts make them prohibitively expensive. `quantmsdiann`'s container-pinned, version-parametrised execution makes this kind of audit a single-command operation.

Several limitations are worth noting. First, `quantmsdiann` does not currently include batch correction, cell-type assignment, or differential-abundance modules; these are best handled in downstream R or Python packages such as `scp` [20], which can read the QPX or MSstats outputs directly. Second, library generation for plexDIA channel deconvolution remains internal to DIA-NN, and `quantmsdiann` inherits any limitations of that step. Third, the QPX format will continue to evolve as community feedback accumulates; backward compatibility is maintained through schema versioning. Fourth, we focused on bulk DIA cell-line reanalyses; plexDIA, TMT-DIA, and single-cell DIA demonstrations are part of the development roadmap but are not evaluated here.

Finally, we note that `quantmsdiann` is complementary to, rather than competing with, existing single-cell DIA workflows. The benchmarking work of Wang *et al.* [4] compared multiple search engines and rescoring strategies; `quantmsdiann` adopts the broadly competitive DIA-NN engine and focuses on the orthogonal problem of orchestrating that engine reproducibly at scale. Similarly, earlier nf-core DIA efforts such as DIAproteomics [21] and, more recently, the MHCquant2 pipeline [22] address analogous orchestration problems for DIA and immunopeptidomics. Together, these contributions extend the nf-core proteomics ecosystem to cover the major MS-based proteomics modalities behind a common SDRF-driven interface.

Declarations

Ethics approval and consent to participate

Not applicable. All datasets reanalysed in this study were obtained from public proteomics repositories with ethics approval and informed consent from

the original depositing studies.

Consent for publication

Not applicable.

Data availability

All datasets reanalysed in this study are publicly available through ProteoMeXchange. Cell-line bulk DIA reanalyses: PRIDE PXD003539 (Guo 2019, NCI-60), PXD004701 (Sun 2023), PXD030304 (ProCan-DepMapSanger), PXD017199 (Tognetti 2021), PXD041421 (Wang 2023). Scalability experiment: PRIDE PXD071075 (Wang 2025). ProteoBench benchmarks: PXD049412 (single-cell, Module 9), PXD062685 (diaPASEF, Module 5), PXD070049 (ZenoTOF, Module 10), and Module 7 (Astral) datasets. The full list of PRIDE accessions used in the pan-cohort of DIA cell-line reanalyses is provided in Additional file 1: Table S1. Per-process runtime and resource envelopes (Figs. S1, S2) and detailed ProteoBench supplementary panels (Figs. S3, S4) are provided as Additional file 1.

Code availability

quantmsdiann is available at <https://github.com/bigbio/quantmsdiann> under the MIT license; documentation is hosted at <https://quantmsdiann.quantms.org/>. The pipeline is implemented in Nextflow DSL2 ($\geq 25.10.4$) and packages each step as a versioned Docker, Singularity, or Apptainer container image. Pipeline version v2.1.0 (release codename “Sao Pablo”, 2026-05-08) was used for the analyses reported in this paper. Analysis scripts that generated the figures in this manuscript are available at <https://github.com/bigbio/quantmsdiann-paper> (this repository).

Competing interests

[VD is the original author of DIA-NN. Other authors declare no competing interests.]

Funding

[Grant numbers and funding sources for each author.]

Authors' contributions

[QXY, YS and CD developed the pipeline. AL and HW implemented the SDRF parsing and QPX export. JN and FC contributed the spatial-proteomics datasets and analysis. VD provided DIA-NN expertise and guidance on cross-version reproducibility. MB and YPR conceived and supervised the project. QXY and YPR wrote the manuscript with input from all authors. All authors read and approved the final manuscript.]

Acknowledgements

[We thank the Brunner, Derks, Leduc, Mund, Wang, and Coscia teams for depositing their data publicly. We thank the nf-core community and the OpenMS and DIA-NN developer teams for their ongoing work.]

References

- [1] V. Demichev, C. B. Messner, S. I. Vernardis, K. S. Lilley, M. Ralser, DIA-NN: neural networks and interference correction enable deep proteome coverage in high throughput, *Nat Methods* 17 (2020) 41–44. doi:10.1038/s41592-019-0638-x.
- [2] A.-D. Brunner, M. Thielert, C. Vasilopoulou, C. Ammar, F. Coscia, A. Mund, et al., Ultra-high sensitivity mass spectrometry quantifies single-cell proteome changes upon perturbation, *Mol Syst Biol* 18 (2022) e10798. doi:10.15252/msb.202110798.
- [3] F. Meier, A.-D. Brunner, M. Frank, A. Ha, I. Bludau, E. Voytik, et al., diaPASEF: parallel accumulation-serial fragmentation combined with data-independent acquisition, *Nat Methods* 17 (2020) 1229–1236. doi:10.1038/s41592-020-00998-0.
- [4] J. Wang, Y. Huang, F. Lu, Q. Xu, Z. Yang, Y. Jiang, et al., Benchmarking informatics workflows for data-independent acquisition single-cell proteomics, *Nat Commun* 16 (2025). doi:10.1038/s41467-025-65174-4.
- [5] E. W. Deutsch, N. Bandeira, Y. Perez-Riverol, V. Sharma, J. J. Carver, L. Mendoza, D. J. Kundu, C. Bandla, S. Kamatchinathan, S. Hewapathirana, et al., The ProteomeXchange consortium in 2026: making proteomics data FAIR, *Nucleic Acids Res* 54 (D1) (2026) D459–D469. doi:10.1093/nar/gkaf1146.

- [6] C. Dai, J. Pfeuffer, H. Wang, P. Zheng, L. Käll, T. Sachsenberg, et al., quantms: a cloud-based pipeline for quantitative proteomics enables the reanalysis of public proteomics data, *Nat Methods* 21 (2024) 1603–1607. doi:10.1038/s41592-024-02343-1.
- [7] C. Dai, A. Füllgrabe, J. Pfeuffer, E. M. Solovyeva, J. Deng, P. Moreno, et al., A proteomics sample metadata representation for multiomics integration and big data analysis, *Nat Commun* 12 (2021) 5854. doi:10.1038/s41467-021-26111-3.
- [8] P. A. Ewels, A. Peltzer, S. Fillinger, H. Patel, J. Alneberg, A. Wilm, et al., The nf-core framework for community-curated bioinformatics pipelines, *Nat Biotechnol* 38 (2020) 276–278. doi:10.1038/s41587-020-0439-x.
- [9] R. Devreese, C. Jachmann, B. Van Puyvelde, H. Webel, W. Bittremieux, C. Chiva, et al., ProteoBench: the community-curated platform for comparing proteomics data analysis workflows, *bioRxivPreprint* (2025). doi:10.64898/2025.12.09.692895.
- [10] N. Hulstaert, J. Shofstahl, T. Sachsenberg, M. Walzer, H. Barsnes, L. Martens, et al., ThermoRawFileParser: modular, scalable, and cross-platform RAW file conversion, *J Proteome Res* 19 (2020) 537–542. doi:10.1021/acs.jproteome.9b00328.
- [11] S. Kamatchinathan, S. Hewapathirana, C. Bandla, S. Insua, J. A. Vizcaíno, Y. Perez-Riverol, pridepy: a Python package to download and search data from PRIDE database, *J Open Source Softw* 10 (107) (2025) 7563. doi:10.21105/joss.07563.
- [12] UniProt Consortium, UniProt: the universal protein knowledgebase in 2025, *Nucleic Acids Res* 53 (2025) D609–D617. doi:10.1093/nar/gkae1010.
- [13] Q.-X. Yue, C. Dai, S. Kamatchinathan, A. Larrea, Y. Perez-Riverol, et al., pmultiqc: an open-source, lightweight, and metadata-oriented QC reporting library for MS proteomics, *bioRxivPreprint* (2025). doi:10.1101/2025.11.02.685980.
- [14] P. Ewels, M. Magnusson, S. Lundin, M. Käller, MultiQC: summarize analysis results for multiple tools and samples in a single report, *Bioinformatics* 32 (2016) 3047–3048. doi:10.1093/bioinformatics/btw354.

- [15] I. Virshup, D. Bredikhin, L. Heumos, G. Palla, G. Sturm, A. Gayoso, et al., The scverse project provides a computational ecosystem for single-cell omics data analysis, *Nat Biotechnol* 41 (2023) 604–606. doi:10.1038/s41587-023-01733-8.
- [16] J. Derks, A. Leduc, G. Wallmann, R. G. Huffman, M. Willetts, S. Khan, et al., Increasing the throughput of sensitive proteomics by plexdia, *Nat Biotechnol* 41 (2023) 50–59. doi:10.1038/s41587-022-01389-w.
- [17] A. Leduc, L. Khoury, J. Cantlon, S. Khan, N. Slavov, Massively parallel sample preparation for multiplexed single-cell proteomics using npop, *Nat Protoc* 19 (2024) 3750–3776. doi:10.1038/s41596-024-01033-8.
- [18] A. Mund, F. Coscia, A. Kriston, R. Hollandi, F. Kovács, A.-D. Brunner, et al., Deep visual proteomics defines single-cell identity and heterogeneity in primary human cancers, *Nat Biotechnol* 40 (2022) 1231–1240. doi:10.1038/s41587-022-01302-5.
- [19] A. Makhmut, D. Qin, S. Fritzsche, J. Nimo, J. König, F. Coscia, A framework for ultra-low-input spatial tissue proteomics, *Cell Syst* 14 (2023) 1002–1014.e5. doi:10.1016/j.cels.2023.10.003.
- [20] C. Vanderaa, L. Gatto, The current state of single-cell proteomics data analysis, *Curr Protoc* 3 (2023) e658. doi:10.1002/cpz1.658.
- [21] L. Bichmann, S. Gupta, G. Rosenberger, L. Kuchenbecker, T. Sachsenberg, P. Ewels, et al., DIAproteomics: a multifunctional data analysis pipeline for data-independent acquisition proteomics and peptidomics, *J Proteome Res* 20 (2021) 3758–3766. doi:10.1021/acs.jproteome.1c00123.
- [22] J. Scheid, S. Lemke, N. Hoenisch-Gravel, A. Dengler, T. Sachsenberg, A. Declercq, et al., MHCquant2 refines immunopeptidomics tumor antigen discovery, *Genome Biol* 26 (2025) 290. doi:10.1186/s13059-025-03763-8.